

PPC Pipeline Supplementary Documentation - Version 2

Johannes Nelson, Matteo Ramina

2025-08-05

Overview

In addition to the documentation within the scripts in the form of comments, this file provides a higher level overview of each script, highlighting key points for functionality, as well as any key assumptions being made.

Extract__Main__Data.R

General

This script's purpose is to take the data out of Kobo, process it and clean it, and produce organized, easily-legible files for further analysis.

Video Link

[Click here for Extraction Overview Video](#)

Inputs/Requirements

Since this is the first stage, there are no local inputs. However, the user will be prompted for the Kobo username and password, so these need to be handy and that user's permissions need to be such that they can extract data in this way.

Outputs

The outputs of this script are files for submission data, tree data, photo data, and geolocation data. They will be stored in a folder called 'Main_Data'.

Assumptions/Notes

There are some assumptions hard-coded into this script.

- The script assumes that the name mappings in the 'Kobo_Key' document are accurate.
- The script assumes naming conventions within Kobo fields does not change with future form versions (e.g. it searches for geolocation data by looking for the strings 'Corner' and 'Centroid'. If these fields are changed in future form versions, it will not be able to find them. The same goes for key data fields like tree count, species, etc.) Here is a coded example:

```
names(df)[grep("number", names(df), ignore.case = TRUE)] <- "Tree_Count"
names(df)[grep("species", names(df), ignore.case = TRUE)] <- "Species"
names(df)[grep("type", names(df), ignore.case = TRUE)] <- "Tree_Type"
```

This chunk of code conforms column labels by searching for key words within the column labels as they are in Kobo. If these change, this will not find them. Also, if a new column is later added with the word number, species, or type as a part of its name, this will try to label that column with the uniform labels incorrectly.

- **Removal of empty columns:** There are a number of entirely empty columns in the dataset that mostly correspond to fields that are not inputs but descriptions of parts of the Kobo form. The assumption here is that if there is not a single piece of information from any of the submissions, it is likely one of these technical fields.
- This script labels key information that is used for analysis like monitoring plot sizes (which are summed to get total area monitored) and size class of trees. Both of these labels are created using logic that considers responses in the forms (e.g. If the SiteSize answer is “No” it is a 10x10 plot.) Trees marked as planted at Y0 are given their own size class of ‘<10cm planted’, since they were ending up in many other tables.

Troubleshooting

- **‘robotoolbox’ R package issue:** At the time of creation, there was a bug in this package, and I needed the creator to go in and fix it, which he was kind enough to do. However, the fix is currently only applied to the development version, so installing this package in the conventional way will not install the version we need. If you get error that looks something like this:

Error in rbindlist(drop_nulls(survey), fill = TRUE): Column 7 of item 3 is length 3 inconsistent with column 1 which is length 5. Only length-1 columns are recycled.

Then you have the older version. Try this:

```
remove.packages("robotoolbox")
install.packages("remotes")
remotes::install_gitlab("dickoa/robotoolbox")
```

- **JSON error:** This happened rarely, but if you encounter an error that talks about JSON format/parsing, wait and try to run the script again. It has to do with a faulty connection/communication step between your machine and Kobo.
- **Table changes within Kobo:** While unlikely at this point in data collection, there is ever-present potential that Kobo will confusingly store things in opaque and mysterious ways that are slightly different than they were at the time I wrote the script. I have included some outputs in the console that should be able to help you see if this is happening, such as the number of expected and retrieved tables. Printed statements are also created at each step of the script, so if an error does appear, you can locate which module created the error easily and address the issue.

Extract_Brazil_Data.R

General

This does the same thing as the above script, and everything above pertains to this one as well. I will talk only about the things that are unique to this script here.

Because of the relatively small amount of submissions present in the database at the time of creating the script, I feel less certain that the submissions (including my test submissions) cover all bases of what might get put into the database in the future. Since the main data is full of strange tables, I assume that with time, some of these peculiarities may appear in the Brazil dataset as well. If this happens, the script will need to be altered to accommodate the changes.

Video

[Click here for Brazil Extraction Video](#)

Assumptions/Notes

- The Brazil Kobo form was designed differently than the primary form. This was to accommodate collection of different data. There are also differences, however, even in the areas that overlap. The video will go into these differences in more detail and highlight areas of potential concern.

Extract_Colombia_Data.R

General

The purpose of this script is to extract the XLSX dataset from the paper form survey conducted in Colombia, and adapt it to a format that facilitates its integration into Kobo's server.

Inputs/Requirements

- “26022025 BD Monitoreo Musesi Parcelas Permanentes RV Final.xlsx”: the digitized version of the data collected in Colombia. This file contains the data to be pre-processed and uploaded.
- “Forms_mapping.xlsx”: a file that links some columns from the Colombian dataset to the corresponding fields in the main Kobo form. It is used to facilitate the selection and renaming of relevant variables.

These inputs should be placed in the “Colombia_Data” sub-folder. The digitized dataset is not included in the GitHub repository.

Outputs

- “main_colombia_data.json”, a JSON file containing the Colombian data, harmonized to match the structure of the main Kobo form for seamless upload.

Assumptions/Notes

The Colombia paper form features some differences compared to the main Kobo form:

- Tree-level data (height, diameter at breast height (DBH), and more) is recorded for each tree within the plots, yielding a dataset that is structured in long (vertical) form. Each observation is uniquely identified by keys representing the date of data collection, the plot ID, the tree ID, and the tree type (planted or already present). This structure is in stark contrast to the main Kobo form's dataset, which is implicitly in wide format, with each observation uniquely identified by the date of data collection and plot ID, and tree information (species and counts) reported horizontally. Moreover, in the Colombia data, nested plots are also reported on separate rows from their main plots and follow the same long format.

- Data for main restoration plot is reported not only from 30x30m plots, as expected, but also from plots of size 10x10m. This departs from the main Kobo form’s sampling approach, which requires restoration data from main plots to be collected exclusively from 30x30m plots (or from 3x3m plots if site size is too small not accommodate a 30x30m plot).
- Information on tree DBH is missing for trees not planted by the project. Due to the Colombia dataset structure and survey logic, this information is critical (DBH data is essential for the fields **Resampling1**, **TreesPresent**, **Resampling2**, and **LittleTreesPresent**). For instance, a tree from a 30x30m plot should be collected in one part of the survey if its DBH is less than 1cm, and in another if its DBH is equal to or greater than 10cm. Without DBH data, it is unclear where that tree’s count should be assigned.
- Resampling was not always conducted according to the instructions in the main Kobo form. For example, some 30x30m plots were not resampled even though the recorded tree data showed no trees with DBH equal to or greater than 10cm. A similar situation occurred for 3x3m nested plots, which were not resampled despite not containing any trees with DBH between 1cm (inclusive) and 10cm (exclusive).
- Partially related to the above, there are instances where trees were recorded even though they should not have been. For example, in a main plot with no trees having DBH equal to or more than 10cm, but with a nested subplot containing trees between 1cm and 10cm DBH, trees in that DBH range were also recorded for the main plot. According to the survey logic, such trees should only be recorded in 3x3m plots. Another example includes planted trees with DBH <1cm being recorded in nested plots, although they should only be recorded in 30x30m plots.

These departures and discrepancies from the main Kobo form required the following assumptions and adjustments:

- The script restructures the data to reflect the implicit wide format required. In doing so, dataframes are created to mimic the spreadsheet obtained when downloading the main Kobo form dataset from Kobo’s server. This results in the **tree_monitoring_sheet**, which contains the data in wide format, along with repeat group sheets for tree counts (**PlantedTrees3**, **group_an2yk58**, **group_ka7vj63**, **group_oq8nt56**, **group_q16qg69**, and **group_wj4lz16**). These dataframes are then joined to form the JSON file used for upload.
- Main restoration plots measuring 10x10m are treated as if they were 30x30m. Tree counts are adjusted via linear extrapolation as follows:

$$\text{Extrapolated Count} = 10\text{x}10\text{m Count} \cdot 9$$

That is, tree counts from 10x10m plots are multiplied by 9. Counts affected by this transformation are flagged in the corresponding note field. The resulting range of extrapolated tree counts is shown in Figure 1. This transformation introduces outliers, as Figure 1 illustrates (e.g., in **group_an2yk58**, the range of extrapolated counts is twice that of 30x30m plots: 35 vs. 14). This needs to be addressed during the analysis stage.

- Upon confirmation of the CI Colombia Team, the missing DBH values for trees that were not planted have been replaced as follows:
 - Already present trees in main plots (30x30m or 10x10m) with missing DBH are assigned a **d1** value of 101 (i.e., 10.1 cm).
 - Already present trees in nested plots with missing DBH are assigned a **d1** value of 50 (i.e., 5 cm).
- Some values for **Resampling1**, **TreesPresent**, **Resampling2**, and **LittleTreesPresent** are adjusted due to the missing DBH data and incorrect resampling recordings described above. Given that **d1** values are recorded in millimeters, the adjustments are as follows:
 - For main plots with **Resampling1** less than 2 but no trees with DBH equal to or above 10cm:
 - * **Resampling1** is updated to 2.
 - * **TreesPresent** is changed to “No”.

- For nested plots of control plots:
 - * **Resampling2** is set to 0. This is justified by the fact that, although DBH measurements are not recorded, it has been confirmed that trees meet the required DBH value, so no resampling is needed.
- For nested plots of restoration plots that have **Resampling2** less than 2 and no trees with DBH equal to or above 1cm and less than 10cm:
 - * **Resampling2** is updated to 2.
 - * **LittleTreesPresent** is changed to “No”.
- Trees that should not have been recorded include two types of planted trees:
 - Planted trees in main plots with DBH in the range $1\text{cm} \leq \text{DBH} < 10\text{cm}$, where the corresponding nested plot also contains planted trees within the same DBH range. In such cases, the trees in the main plot should not have been recorded (they should only appear in **group_ql6qg69** if no tree with that DBH range was found in the nested plot). Nevertheless, to avoid losing their counts, these trees are included in the **PlantedTrees3** group. They are flagged in the note field.
 - Planted trees in nested plots with DBH less than 1cm. Trees with such DBH should only be recorded in the main plots and would typically be included in the **PlantedTrees3** group. Nevertheless, to avoid losing their counts, these trees are included in the **PlantedTrees3** group. They are likewise flagged in the note field.

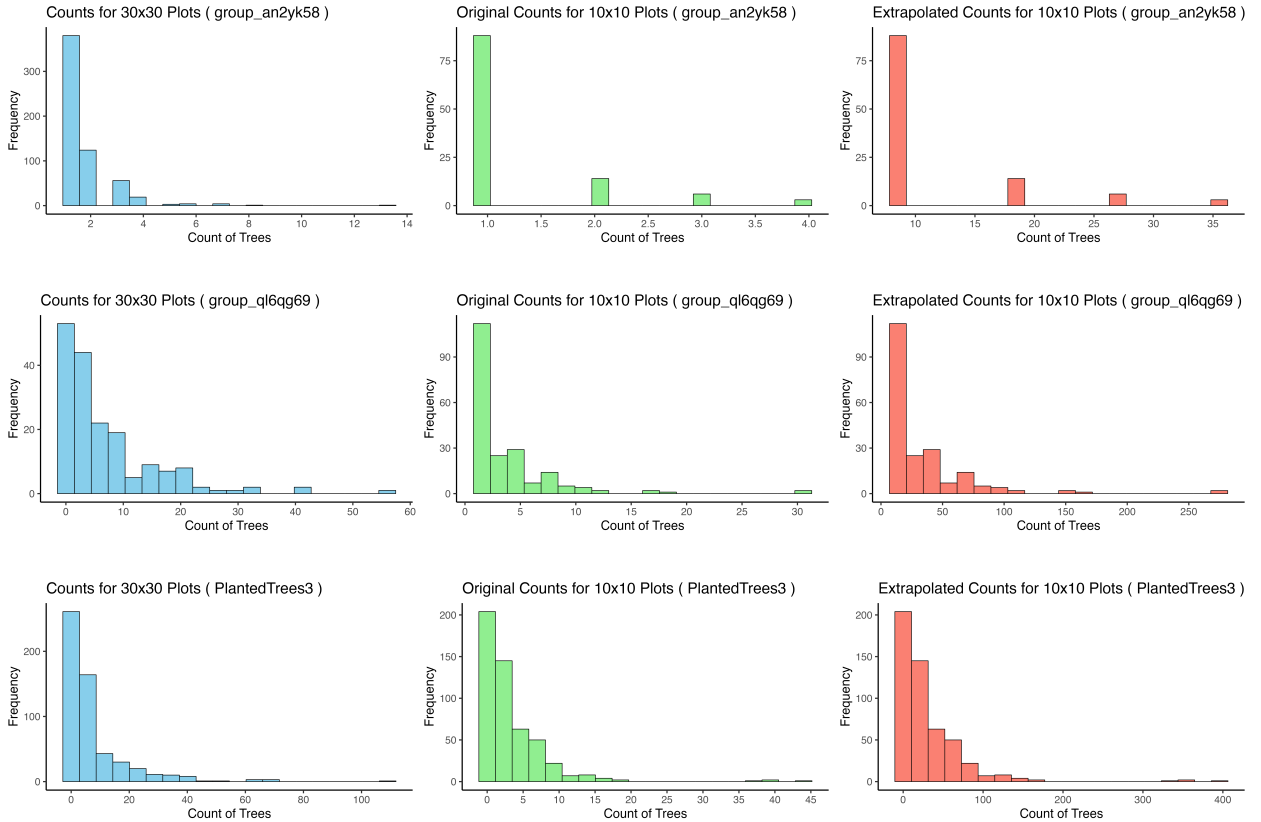


Figure 1: Comparison of tree count ranges across groups and plot sizes.

- The script assumes that key column names in the Excel data will remain consistent across different rounds of data collection. This is because the renaming of these columns is hardcoded at the beginning

of the script, under the `# Rename columns in colombia_data_mod` section. To make future modifications easier, any elements you might need to change, such as DBH column names or file names, are highlighted with a `# CHANGE HERE` indicator on their respective lines.

build_payload.R

General

This script defines a utility function used within “upload_data.R” to generate the payload required for submitting data to Kobo’s server.

upload_data.R

General

This script handles the automated uploading of the data from the Colombia paper form to KoboToolbox.

Inputs/Requirements

- “main_colombia_data.json”: a JSON file containing the Colombian data, harmonized to match the structure of the main Kobo form for seamless upload. It can be found in the “Colombia_Data” subfolder.
- “build_payload.R”: an R script that creates a nested payload formatted for the KoboToolbox API by taking a single data record and a form ID, extracting the UUID, and organizing the submission and metadata into the required structure.
- “.Renviron”, a (hidden) file that is safely stored in your local machine to store your API key. See more details below.
- Two pieces of information are required to run the script:
 - The API key: a personal token that authenticate the user when submitting a POST request.
 - The form ID: a code that uniquely identifies the form that the user wants to upload data for. The user must have access to the form.

Outputs

- “upload_log.txt”: a text file logging the upload process. It is used to verify whether all observations were uploaded on the server.

Moreover, As a result of running this script, the data is uploaded on Kobo’s server.

Assumptions/Notes

The script works in conjunction with build_payload.R, iterating through each observation in the JSON file, generating the corresponding payload, submitting it to the server, and continuing to the next.

To make future modifications easier, any elements you might need to change, such as DBH column names or file names, are highlighted with a `# CHANGE HERE` indicator on their respective lines.

Creation of .Renviron To run the analysis script, you must first securely store your personal KoboToolbox API key by creating a (hidden) file named “.Renviron” that R automatically checks for at the start of a session (for more information, see <https://rstats.wtf/r-startup.html#renviro>). This is achieved as follows:

1. Create the “.Renviron” file and place it in the project’s root folder (PPC_Data_Analysis_Pipeline/). For instructions on creating dotfiles, see this guide for Windows and macOS.
2. Add your API key by opening “.Renviron” in a text editor and add the following line, using your real key: `KOBO_API_TOKEN="paste_your_API_key_here"`
3. Save the file and, if you had an active R session, **restart your R session**. This step is mandatory, as R only reads this file on startup.

The following error could appear when initializing the API key:

Error: API token not found. Please set the ‘KOBO_API_TOKEN’ environment variable.

Solutions to this error are:

- Restart R: This is the most common fix. You must restart your R session after saving the file.
- Check for Typos: Ensure the variable name is exactly “KOBO_API_TOKEN” and that the syntax of the line is correct.
- Check File Name and Location: The file must be named “.Renviron” (not “. Renviron.txt”) and be in the main project folder.
- Manual Override: If nothing else works, force the script to read the file by adding this line to the very top of your R script, using the full path to your file: `readRenviron("/path/to/your/project/PPC_Data_Analysis_Pipeline/.Renviron")`

Correct_Species_Names.R

General

This script uses the ‘taxize’ package to facilitate connection the the Global Names Resolver while attempting to map species names present in the Kobo database with accepted species names in known scientific databases.

Video Link Click here for Species Name Correction Video

Inputs/Requirements

- It will take the extracted data from the previous scripts as input—primarily the file with “Tree_Data_Uncorrected” within its name.
- If one exists, which it always should since the repository contains one at the time of hand-over, a ‘Taxonomic_Corrections’ file in the ‘Species_Data’ folder will also be used to automatically make corrections that are already known. The script will work without this file, but it will take a very long time to run (20+ minutes just for the automated part—much longer for manual validation).

Outputs

- Tree_Data_Corrected_YYYY-MM-DD_HHMM.csv: tree data with corrected names.
- Taxonomic_Corrections_YYY-MM-DD_HHMM.csv: updated taxonomic corrections.
- Optionally, at the end the user can be prompted to get a file showing the species that need review.

Assumptions

- Extraction scripts have been run
- Two databases are used for name matching: Tropicos - Missouri Botanical Garden and The International Plant Names Index. Others are available if these are not desired.

Troubleshooting

- Runtime: as mentioned, there is a part of the script that takes a while to run, so if you think your R session is frozen, it probably has not.
- Manual validation: This can take a long time, since the script will always show you all the species that it was not able to resolve, and prompt you for the correct name. Hundreds of these are likely common names. If you've corrected all the species names recently and are running the script again, it will still show you all these common names one by one, which can be onerous. I included a 'hidden' option, where if you type the word 'save' as a response, it will move forward and assume you know no corrections for those names.

Add_Family_Names.R

General

This script takes the cleaned and corrected species names and attempts to find family names for them. It does this using the 'rgbif' package to interact with the GBIF database. There is also 'hidden' functionality to use the 'taxize' package to pull from the ncbi database. Results were only slightly different, and since the GBIF method was much faster and required no API token set up (which would have been an obnoxious step that each user would have to do before running the first time), I choose the GBIF route. If there is interest in the other way, I can facilitate that too.

Video Link [Click here for Add Family Names Video](#)

Inputs

- Tree_Data_Corrected_YYYY-MM-DD_HHMM.csv: the output from the previous script with corrected species names.
- Family_Names_YYYY-MM-DD_HHMM.csv: the most up to date file with known family names for each species.

Outputs

- Final_Tree_Data_YYY-MM-DD_HHMM.csv: the final tree data file (before passing to analysis), with correct species names and family names added.
- Family_Names_YYYY-MM-DD_HHMM.csv: An up to date family names file.

Assumptions/Notes

- No major assumptions other than the reliability of the GBIF database.

Troubleshooting

- This script will ask the user if they want to manually add family names for those species which could not be automatically matched. If choosing yes, the user will then have to go through each unresolved species name and look up a family name for it. Only run this if you have time. Unlike species names, I did not have time to create a manually up-to-date family names file, so this will require some effort the first time it is run.

Analyze_Data.R

General

This script takes the final tree data, with corrected species names and added family names, and it uses logic based on table names, size class, and resampling to first scale the tree counts to a uniform 30x30m (900m²) size, and then further scales the counts to the size of the total site (if that size information was available.)

Video Link

[Click here for Analyze Data Video](#)

Inputs

- Final_Tree_Data_YYYY-MM-DD_HHMM.csv: the output from the previous script with corrected species names.
- site_size_data.csv: this is a manually saved and located csv that contains site size information.
- Main_Data_YYYY-MM-DD.csv: the main data file created during extraction

Outputs

- Baseline Reports: A folder containing baseline counts of different tree types and different size classes at the site and the species level.
- Indicator Reports: A folder containing files showing planted tree survival rates, trees restored, and trees naturally regenerated—each indicator broken down to the site level and the species level.

Assumptions/Notes

- Trees at Y0 that were marked as planted but that were NOT in the planted tables are ignored.
- Trees at Y0 that were marked as naturally regenerating are included in the counts for ‘already present’ in both the baseline reports and the trees restored reports (where the Y0 count is subtracted)
- There were still some unresolved questions about how to deal with different tree types during later timeframes (e.g. What about trees marked as ‘present’ in later time frames? Are the monitors supposed to remember what trees were already there? What about naturally regenerating?) Questions like this will likely need to be addressed as the data comes in.
- Site IDs that are shared within countries and organizations are actually the same sites.

Troubleshooting

The outputs will appear short or blank (depending on which report), in the absence of data from various timeframes. The script has been tested with randomly generated data for Y25 and Y5, and it runs. Until organizations add data from these timeframes, though, the outputs will not be complete.

IMP_Invasive_Species_Scanner.R

General

This script scans the IMP planted tree and seed data for potentially invasive species.

Video Link

[Click here for IMP Invasive Scanner Video](#)

Inputs

- `ppc_export_xxx...csv`: The csv export with planting data from the IMP that is manually exported from the IMP and placed in the `IMP_Data` folder of this repo

Outputs

- `Invasive_Species_Data_YYYY-MM-DD_HHMM.csv`: a record of all species passed to the script for processing.
- `Invasive_Species_Report_YYY-MM-DD_HHMM.csv`: The report that contains all flagged species, organizations, countries, and relevant info from the GISD database.

Assumptions/Notes

- `ppc_export` file exists (this has to be manually downloaded and placed in the `IMP_Data` folder)
- Common names are not going to return any results. There is a built-in species correction module here as a part of the preprocessing step. There are still common names present in the IMP data, and these will not get flagged even if they are invasive.

Troubleshooting

No known issues.

IMP_Native_Alien_Classification.R

General

This script enables the labeling of species-country pairs as Native, Alien, or Needing further review. If more options are desired, they can be pretty easily coded in. This is a workflow that requires user input, and it can take a long time. If you don't want to do them all at once, you can type 'save' into the response, and this will break the loop and save what you've done. The subsequent run will then check the manually reviewed file and be able to pick up where you left off.

Video Link

[Click here for IMP Native Alien Classification Video](#)

Inputs

- IMP_planted_trees_YYYY-MM-DD_HHMM.csv: preprocessed IMP data output by the Invasive Scanner script
- IMP_planted_seeds_YYYY-MM-DD_HHMM.csv: preprocessed IMP data output by the Invasive Scanner script
- GBIF_dataset_keys.csv: file with GRIIS datasets and their relevant keys for lookup in the GBIF database
- Taxonomic_Corrections_YYYY-MM-DD_HHMM.csv: the most recent species corrections.

Outputs

- Checklist_Scan_Results_YYYY-MM-DD_HHMM.csv: results of checks against GRIIS databases. Though supplementary. While these results are mainly displayed during the manual validation process, this output could also be used as a potential introduced species flagged by checking for what species had matches.
- Taxonomic_Corrections_YYYY-MM-DD_HHMM.csv: an updated species correction file
- Manually_Reviewed_Planting_Data_YYYY-MM-DD_HHMM.csv: a dataset containing the reviewed and labelled data after manual validation.

Assumptions/Notes

- GBIF occurrence data is used for mapping—assumption here is about accuracy and validity

Troubleshooting

- Wikipedia is source for relevant info—sometimes the query might grab an irrelevant page